

SARSA vs Q-learning

SARSA Algorithm

ϵ -greedy policy improvement done for TD methods

-
- 1: Set initial ϵ -greedy policy π , $t = 0$, initial state $s_t = s_0$
 - 2: Take $a_t \sim \pi(s_t)$ // Sample action from policy
 - 3: Observe (r_t, s_{t+1})
 - 4: **loop**
 - 5: Take action $a_{t+1} \sim \pi(s_{t+1})$
 - 6: Observe (r_{t+1}, s_{t+2})
 - 7: $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$
 - 8: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random
 - 9: $t = t + 1$
 - 10: **end loop**
- Loop Q value update for every time-step : based on TD methods
- E-greedy policy update?
-

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S

Choose A from S using policy derived from Q (e.g., ε -greedy)

Loop for each step of episode:

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal

Action chosen by ε -greedy
Action taken for next state already decided

```
1: Set initial  $\epsilon$ -greedy policy  $\pi$ ,  $t = 0$ , initial state  $s_t = s_0$ 
2: Take  $a_t \sim \pi(s_t)$  // Sample action from policy
3: Observe  $(r_t, s_{t+1})$ 
4: loop
5:   Take action  $a_{t+1} \sim \pi(s_{t+1})$ 
6:   Observe  $(r_{t+1}, s_{t+2})$ 
7:    $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$ 
8:    $\pi(s_t) = \arg \max_a Q(s_t, a)$  w.prob  $1 - \epsilon$ , else random
9:    $t = t + 1$ 
10: end loop
```

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

```

1: Set initial  $\epsilon$ -greedy policy  $\pi$ ,  $t = 0$ , initial state  $s_t = s_0$ 
2: Take  $a_t \sim \pi(s_t)$  // Sample action from policy
3: Observe  $(r_t, s_{t+1})$ 
4: loop
5:   Take action  $a_{t+1} \sim \pi(s_{t+1})$ 
6:   Observe  $(r_{t+1}, s_{t+2})$ 
7:    $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$ 
8:    $\pi(s_t) = \arg \max_a Q(s_t, a)$  w.prob  $1 - \epsilon$ , else random
9:    $t = t + 1$ 
10: end loop

```

Loop Q value update for every time-step : based on TD methods

E-greedy policy update?

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\epsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Loop for each step of episode:

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal

Action chosen by e-greedy

Action taken for next state already decided

Difference?

Seems to indicate that actions are taken greedily

```

1: Set initial  $\epsilon$ -greedy policy  $\pi$ ,  $t = 0$ , initial state  $s_t = s_0$ 
2: Take  $a_t \sim \pi(s_t)$  // Sample action from policy
3: Observe  $(r_t, s_{t+1})$ 
4: loop
5:   Take action  $a_{t+1} \sim \pi(s_{t+1})$ 
6:   Observe  $(r_{t+1}, s_{t+2})$ 
7:    $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$ 
8:    $\pi(s_t) = \arg \max_a Q(s_t, a)$  w.prob  $1 - \epsilon$ , else random
9:    $t = t + 1$ 
10: end loop

```

Loop Q value update for every time-step based on TD methods

E-greedy policy update?

No indication of how next action is chosen
But can assume greedy action choosing

But this would mean q-value has been calculated for all actions at state s_t

Sutton
and
Barto

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\epsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Loop for each step of episode:

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal

Action chosen by ϵ -greedy
Action taken for next state already decided

lecture 4

stop

Simple difference between SARSA and Q-learning

SARSA = on-policy

Q-learning = off-policy

SARSA = on-policy

on-policy = on e-greedy policy

on-policy = learn to estimate and evaluate policy from experience obtained from following that policy (purely from experience)

CS234
lecture 4

-
- 1: Set initial ϵ -greedy policy π , $t = 0$, initial state $s_t = s_0$
 - 2: Take $a_t \sim \pi(s_t)$ // Sample action from policy
 - 3: Observe (r_t, s_{t+1})
 - 4: **loop**
 - 5: Take action $a_{t+1} \sim \pi(s_{t+1})$
 - 6: Observe (r_{t+1}, s_{t+2})
 - 7: $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$
 - 8: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random
 - 9: $t = t + 1$
 - 10: **end loop**
-
- Loop Q value update for every time-step : based on TD methods

E-greedy policy update?

Sutton
and
Barto

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\epsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Action chosen by e-greedy
Action taken for next state already decided

SARSA = on-policy

on-policy = on ϵ -greedy policy

on-policy = learn to estimate and evaluate policy from experience obtained from following that policy (purely from experience)

CS234
lecture 4

```
1: Set initial  $\epsilon$ -greedy policy  $\pi$ ,  $t = 0$ , initial state  $s_t = s_0$ 
2: Take  $a_t \sim \pi(s_t)$  // Sample action from policy
3: Observe  $(r_t, s_{t+1})$ 
4: loop
5:   Take action  $a_{t+1} \sim \pi(s_{t+1})$ 
6:   Observe  $(r_{t+1}, s_{t+2})$ 
7:    $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$ 
8:    $\pi(s_t) \leftarrow \arg \max_a Q(s_t, a)$  w. prob  $1 - \epsilon$ , else random
9:    $t = t + 1$ 
10: end loop
```

Loop Q value update for every time-step : based on TD methods

ϵ -greedy policy update?

Seems to be experience following a different policy (trying out all other actions)

Sutton
and
Barto

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\epsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Loop for each step of episode:

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal

Action chosen by ϵ -greedy

Action taken for next state already decided

The policy on the
lefthandside seems
stochastic?

8: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random

When choosing action from this policy

If pi takes a single action = either
argmax or random by epsilon

Then

pi = Deterministic policy

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Action already pre-
determined

Conditions of SARSA:

GLIE
Robbin-Munro

Conditions of SARSA:

GLIE
Robbin-Munro

Definition of GLIE

- All state-action pairs are visited an infinite number of times

$$\lim_{i \rightarrow \infty} N_i(s, a) \rightarrow \infty$$

- Behavior policy converges to greedy policy
 $\lim_{i \rightarrow \infty} \pi(a|s) \rightarrow \arg \max_a Q(s, a)$ with probability 1

Basically being able to
explore all actions

= use method of e-
greedy where e starts
big and then diminishes

All s, a are visited
infinitely often

Conditions of SARSA:

GLIE
Robbin-Munro

② The step-sizes α_t satisfy the Robbins-Munro sequence such that

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

Learning rate not too
big..?

Conditions of SARSA:

GLIE
Robbin-Munro

② The step-sizes α_t satisfy the Robbins-Munro sequence such that

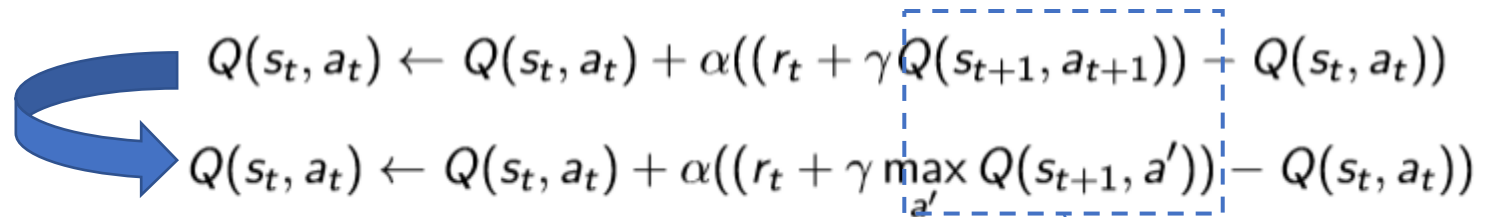
$$\sum_{t=1}^{\infty} \alpha_t = \infty$$
$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

Learning rate not too
big..?

Alpha = 0 : no learning

Alpha = 1 : new experience
complete overtakes

SARSA \rightarrow Q-Learning


$$\begin{aligned} & Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_t + \gamma Q(s_{t+1}, a_{t+1})) - Q(s_t, a_t)) \\ & Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_t + \gamma \max_{a'} Q(s_{t+1}, a')) - Q(s_t, a_t)) \end{aligned}$$

Bootstrapping

-
-
- 1: Initialize $Q(s, a), \forall s \in S, a \in A$ $t = 0$, initial state $s_t = s_0$
 - 2: Set π_b to be ϵ -greedy w.r.t. Q
 - 3: **loop**
 - 4: Take $a_t \sim \pi_b(s_t)$ // Sample action from policy
 - 5: Observe (r_t, s_{t+1})
 - 6: $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma \arg \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$
 - 7: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random
 - 8: $t = t + 1$
 - 9: **end loop**
-

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

Q-learning = off policy

Off-policy = estimate and evaluate policy from following a different policy

-
-
- 1: Initialize $Q(s, a), \forall s \in S, a \in A$ $t = 0$, initial state $s_t = s_0$
 - 2: Set π_b to be ϵ -greedy w.r.t. Q
 - 3: **loop**
 - 4: Take $a_t \sim \pi_b(s_t)$ // Sample action from policy
 - 5: Observe (r_t, s_{t+1})
 - 6: $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma \arg \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$
 - 7: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random
 - 8: $t = t + 1$
 - 9: **end loop**
-

Not necessarily the same policy



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$
 Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$
 Loop for each episode:
 Initialize S
 Loop for each step of episode:
 Choose A from S using policy derived from Q (e.g., ε -greedy)
 Take action A , observe R, S'
 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
 $S \leftarrow S'$
 until S is terminal

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$
 Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$
 Loop for each episode:
 Initialize S
 Choose A from S using policy derived from Q (e.g., ε -greedy)
 Loop for each step of episode:
 Take action A , observe R, S'
 Choose A' from S' using policy derived from Q (e.g., ε -greedy)
 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$
 $S \leftarrow S'; A \leftarrow A'$
 until S is terminal

Not necessarily the same
policy

Difference = random by epsilon?

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$
 Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$
 Loop for each episode:
 Initialize S
 Loop for each step of episode:
 Choose A from S using policy derived from Q (e.g., ε -greedy)
 Take action A , observe R, S'
 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
 $S \leftarrow S'$
 until S is terminal

Not necessarily the same policy

Difference =

No matter what next action chosen for evaluating current state,

Next action will be decided with the updated policy (new and different from that used to evaluate)

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$
 Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$
 Loop for each episode:
 Initialize S
 Choose A from S using policy derived from Q (e.g., ε -greedy)
 Loop for each step of episode:
 Take action A , observe R, S'
 Choose A' from S' using policy derived from Q (e.g., ε -greedy)
 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$
 $S \leftarrow S'; A \leftarrow A'$
 until S is terminal

When evaluated with next action chosen, the chosen next action will be used on next step

Q-learning tends to have better results

For most cases...

But for some cases...

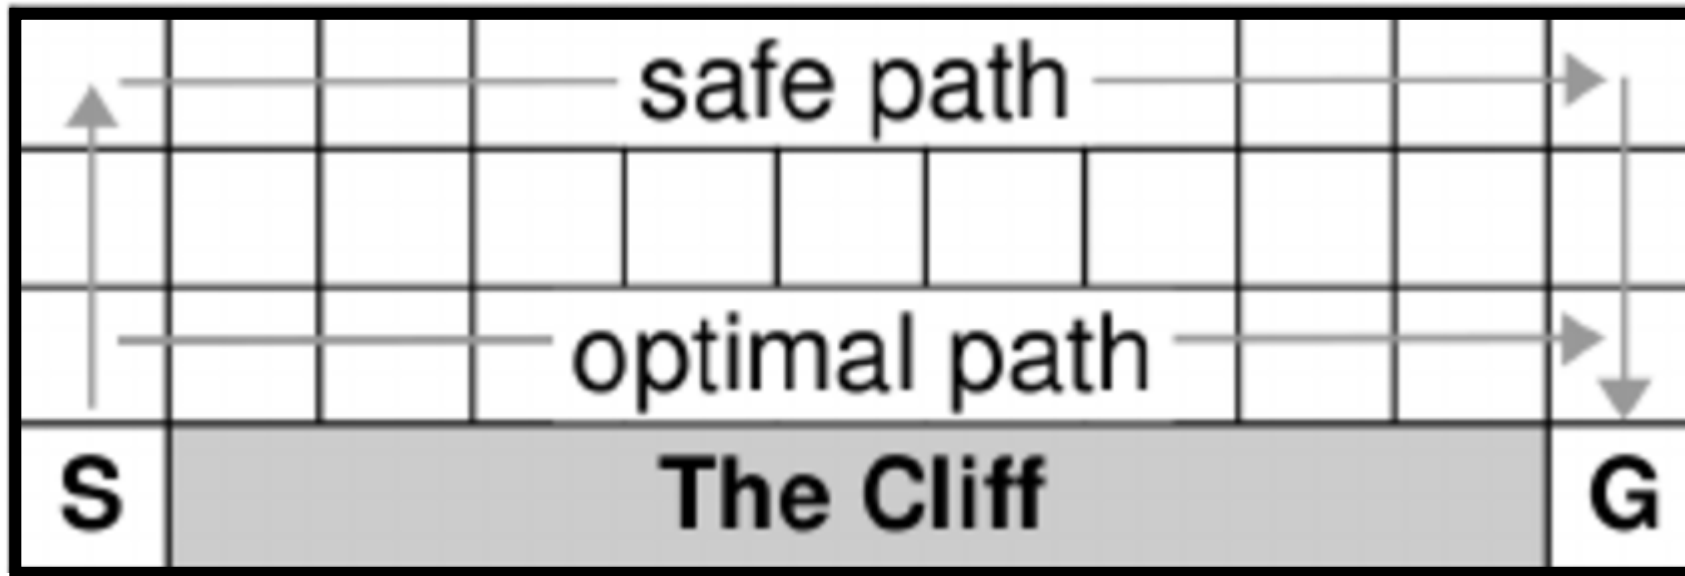
Why SARSA rather than Q-learning?

Sutton and Barto example of cliff-walking

SARSA may do better on early stages

SARSA is realistic

Q-learning is optimistic



Case where there could be a lot of negative outcomes from optimism of using max

Optimism would mean being only interested in the best outcome

Optimism on Cliff for example:

Near cliff is the shortest route to finish = highest possible reward & but forgets that it also has high chance of falling

SARSA Code breakdown

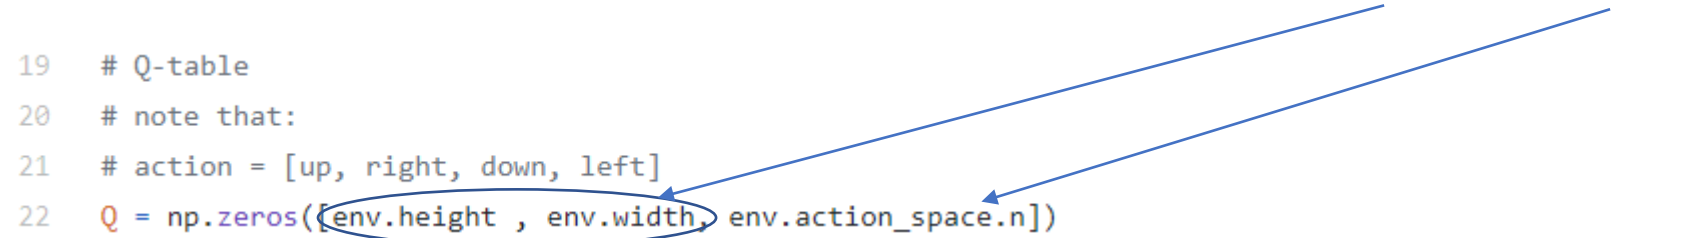
```
9  import gym
10 import gym_gridworlds ← Import package that makes &
11 import numpy as np      registers Cliff environment
12 import matplotlib.pyplot as plt
13 from gym.envs.registration import register
14 import random as pr
15
16
17 env = gym.make('Cliff-v0') ← Create element
```

SARSA Code breakdown

Q value (state-action value) is defined by value that comes from action taken at current state

= preparing Table for all possible states and actions

```
19 # Q-table
20 # note that:
21 # action = [up, right, down, left]
22 Q = np.zeros([env.height, env.width, env.action_space.n])
```



SARSA Code breakdown

```
19 # Q-table
20 # note that:
21 # action = [up, right, down, left]
22 Q = np.zeros([env.height, env.width, env.action_space.n])
```

Current state can be defined by current position (x, y)



SARSA Code breakdown

```
24 # define hyperparameters (epsilon defined as function later)
25 dis = .91
26 num_episodes = 1000
27 lr = 0.01
28 success = 0
29
30 # create lists to contain total rewards and steps per episode = to track how total reward changes
31 rList = []
```

SARSA Code breakdown

```
def reset(self):
    self.S = (3, 0)
    return self.S
```

```

32 for i in range(num_episodes):
33     # Reset environment and get first new ob
34     (y, x) = env.reset()
35     # state = y , x
36     rAll = 0
37     done = False
38
39     e = 1. / ((i // 100) + 1)
40     # e-greedy where epsilon starts big
41     # and becomes small
42     # = more exploration in the beginning and less later on
43     # with larger num_episodes, this might come into better effect
44
45     # note that this is a method to satisfy condition of GLIE

```

R= -1	R= -1	R= -1	R= -1	
R= -1	R= -1	R= -1	R= -1	
R= -1	R= -1	R= -1	R= -1	
S				

R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1
R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1
R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1	R= -1
S	The Cliff (R = -100)								G

SARSA Code breakdown

```
32  for i in range(num_episodes):
33      # Reset environment and get first new observation
34      (y, x) = env.reset()
35      # state = y , x
36      rAll = 0
37      done = False
38
39      e = 1. / ((i // 100) + 1)
40      # e-greedy where epsilon starts big
41      # and becomes small
42      # = more exploration in the beginning and less later on
43      # with larger num_episodes, this might come into better effect
44
45      # note that this is a method to satisfy condition of GLIE
```

Done = reached destination

When done = True, episode
is terminated

As #episode trained
increases,

Epsilon decreases

SARSA Code breakdown

```
47     # e-greedy
48     if np.random.rand(1) < e:
49         action = env.action_space.sample()
50     else:
51         action = np.argmax(Q[y, x, :])
```

SARSA Code breakdown

```
53     # The Q-Table learning algorithm
54     while not done:
55         # Get new state and reward from environment
56         (new_y, new_x), reward, done, _ = env.step(action)
57         # new_state = new_x , new_y
58         # Choose the next action by e greedy
59         if np.random.rand(1) < e:
60             next_action = env.action_space.sample()
61         else:
62             next_action = np.argmax(Q[new_y, new_x, :])
```


SARSA Code breakdown

```
53     # The Q-Table learning algorithm
54     while not done:
55         # Get new state and reward from environment
56         (new_y, new_x), reward, done, _ = env.step(action)
```

```
24     def step(self, action):
25         x, y = self.moves[action]
26         self.S = self.S[0] + x, self.S[1] + y
27
28         self.S = max(0, self.S[0]), max(0, self.S[1])
29         self.S = (min(self.S[0], self.height - 1),
30                  min(self.S[1], self.width - 1))
31
32         if self.S == (self.height - 1, self.width - 1):
33             return self.S, -1, True, {}
34         elif self.S[1] != 0 and self.S[0] == self.height - 1:
35             # the cliff
36             return self.S, -100, True, {}
37         return self.S, -1, False, {}
```

```
14         self.moves = {
15             0: (-1, 0),    # up
16             1: (0, 1),    # right
17             2: (1, 0),    # down
18             3: (0, -1),   # left
19         }
```

SARSA Code breakdown

```
64     # Update Q-Table with new knowledge using learning rate
65     Q[y, x, action] = (Q[y, x, action] + lr * (reward + dis * Q[new_y, new_x, next_action] - Q[y, x, action]))
66
67     rAll += reward
68     y, x = new_y, new_x
69
70     # difference with Q-learning
71     # action is already chosen for updating current state-action value
72     action = next_action
73     # therefore the loop continues by choosing only the next action
74
75     rList.append(rAll)
```

Evaluate using
next_action determined
by the current policy

SARSA Code breakdown

```
64     # Update Q-Table with new knowledge using learning rate
65     Q[y, x, action] = (Q[y, x, action] + lr * (reward + dis * Q[new_y, new_x, next_action] - Q[y, x, action]))
66
67     rAll += reward
68     y, x = new_y, new_x
69
70     # difference with Q-learning
71     # action is already chosen for updating current state-action value
72     action = next_action
73     # therefore the loop continues by choosing only the next action
74
75 rList.append(rAll)
```

Then use the
next_action for
next_state in next
timestep

SARSA Code breakdown

```
64     # Update Q-Table with new knowledge using learning rate
65     Q[y, x, action] = (Q[y, x, action] + lr * (reward + dis * Q[new_y, new_x, next_action] - Q[y, x, action]))
66
67     rAll += reward
68     y, x = new_y, new_x
69
70     # difference with Q-learning
71     # action is already chosen for updating current state-action value
72     action = next_action
73     # therefore the loop continues by choosing only the next action
74
75     rList.append(rAll)
```



Loop until termination

SARSA Code breakdown

```
77 print("Success rate: " + str(sum(rList) / num_episodes))
78 print("Final Q-Table Values")
79 print(Q)
80 print("Q Values near starting position")
81 print("START\n" + str(Q[3,0:1]))
82 print("Above START\n" + str(Q[2,0:2]))
83 print("Above above START\n" + str(Q[1,0:2]))
84 plt.bar(range(len(rList)), rList, color="blue")
85 plt.show()
```

Q-learning Code breakdown

```
31 # The Q-Table learning algorithm
32 while not done:
33
34     if np.random.rand(1) < e:
35         action = env.action_space.sample()
36     else:
37         action = np.argmax(Q[y, x, :])
38
39     # Get new state and reward from environment
40     (new_y, new_x), reward, done, _ = env.step(action)
41     # new_state = new_x , new_y
42
43     # Update Q-Table with new knowledge using learning rate
44     Q[y, x, action] = (Q[y, x, action] + lr * (reward + dis *
45     np.max(Q[new_y, new_x, :]) - Q[y, x, action]))
46
47     rAll += reward
48     y, x = new_y, new_x
49
50 rList.append(rAll)
```

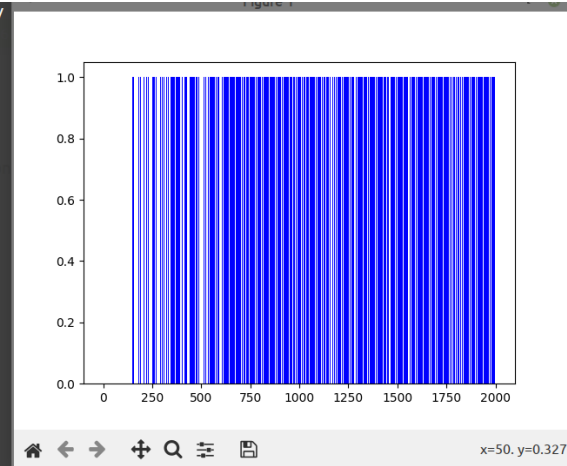
Second difference

Main difference

Q-learning tends to have better results


SARSA

```
laphisboy@laphisboy-900X3K:~/RL_fall/fall_week_1$ python3 just_lake_SARSA.py
Success rate: 0.7905
Final Q-Table Values
[[6.70736658e-02 4.77326352e-01 5.80999368e-03 8.48964303e-02]
 [1.05166556e-01 0.00000000e+00 1.61637981e-07 3.68353915e-04]
 [1.03474654e-02 1.90521207e-07 1.18426275e-11 1.21421652e-06]
 [5.78286553e-07 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [7.22760373e-02 5.54426264e-01 0.00000000e+00 5.51919838e-02]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 7.07570326e-04]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [1.24071121e-01 0.00000000e+00 6.44242082e-01 9.22765749e-02]
 [9.58373591e-02 5.85828924e-02 7.64135946e-01 0.00000000e+00]
 [1.57588989e-01 8.84203921e-01 0.00000000e+00 3.71905143e-05]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [0.00000000e+00 6.08529223e-03 4.09755896e-01 5.92824654e-03]
 [5.26202400e-02 2.62430749e-01 9.9999874e-01 1.93303018e-01]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]]
```



Q-Learning


```
laphisboy@laphisboy-900X3K:~/RL_fall/fall_week_1$ python3 just_lake_q.py
Score over time: 0.81
Final Q-Table Values
[[0.531441 0.59049 0.59049 0.531441 ]
 [0.531441 0. 0.6561 0.59048989]
 [0.59036441 0.729 0.4782969 0.64894266]
 [0.649539 0. 0. 0. ]
 [0.59049 0.6561 0. 0.531441 ]
 [0. 0. 0. 0. ]
 [0. 0.81 0. 0.64829366]
 [0. 0. 0. 0. ]
 [0.6561 0. 0.729 0.59049 ]
 [0.6561 0.81 0.81 0. ]
 [0.72899927 0.9 0. 0.72826444]
 [0. 0. 0. 0. ]
 [0. 0. 0. 0. ]
 [0. 0.81 0.9 0.729 ]
 [0.81 0.9 1. 0.81 ]
 [0. 0. 0. 0. ]]
```



Q-learning tends to have better results


SARSA

```
laphisboy@laphisboy-900X3K:~/RL_fall/fall_week_1$ python3 frozen_lake_SARSA.py
Success rate: 0.0545
Final Q-Table Values
[[1.55087990e-03 6.78733437e-03 1.30792462e-03 9.31839567e-04]
 [1.04827592e-04 3.51441482e-03 2.05297738e-04 4.02499257e-04]
 [3.87858024e-04 7.99769647e-03 5.85235056e-04 5.91762768e-05]
 [1.26320704e-04 1.84382247e-03 5.84039809e-06 3.39594415e-05]
 [1.38777383e-02 1.73185685e-03 2.00942524e-03 6.09587098e-04]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [2.71560631e-02 1.22320626e-10 1.33757456e-03 1.99071938e-05]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [1.78146584e-03 2.79410650e-03 2.46060202e-03 2.91133741e-02]
 [3.38610920e-03 5.83086773e-03 6.36602543e-02 3.06886333e-03]
 [6.92414686e-03 1.09936351e-01 3.59391950e-03 2.88692781e-04]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [3.52018257e-03 1.28705679e-01 1.44792346e-02 1.61855613e-02]
 [3.28335009e-03 3.20059102e-02 4.79574003e-03 3.53480997e-01]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]]
```



Q-Learning

```
laphisboy@laphisboy-900X3K:~/RL_fall/fall_week_1$ python3 frozen_lake_q.py
Score over time: 0.1205
Final Q-Table Values
[[4.30972397e-02 8.68329874e-02 5.16130150e-03 1.39912677e-02]
 [8.41989861e-04 1.11620981e-03 2.78075826e-04 4.70342591e-02]
 [1.04278780e-03 6.54248387e-03 1.70384432e-03 1.56023476e-03]
 [1.34461696e-03 3.82147058e-05 1.81494164e-04 1.63379968e-03]
 [4.92277658e-02 1.55019952e-03 1.01424568e-02 3.88767743e-03]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [5.30290511e-07 6.58845984e-07 8.85605526e-04 2.62413302e-06]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [1.11079066e-03 1.84759192e-02 3.67341763e-02 3.01863580e-01]
 [4.05955922e-04 5.21394469e-01 3.49797214e-02 5.06608931e-02]
 [1.12056118e-01 2.71397368e-06 8.23377276e-06 2.83359708e-02]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [2.11084497e-02 3.40265637e-01 8.54451439e-01 2.59187179e-03]
 [1.73027131e-01 9.89435350e-01 1.48444394e-01 4.63286997e-01]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]]
```

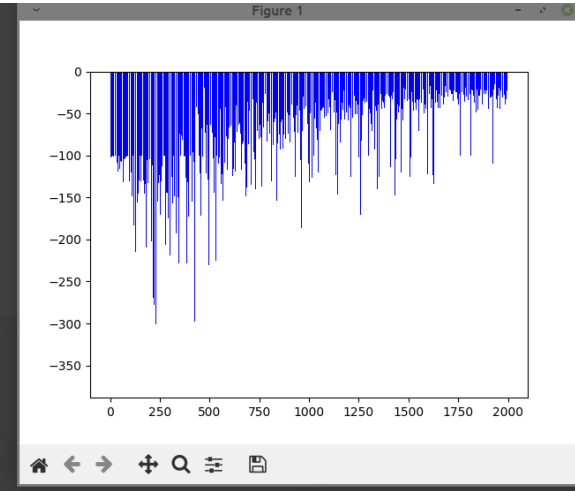


Why SARSA rather than Q-learning?

SARSA

```
laphisboy@laphisboy-900X3K:~/RL_fall/fall_week_1$ python3 cliff_SARSA.py
Success rate: -78.4395
Final Q-Table Values
[[[-6.57224026 -6.56897259 -6.57255013 -6.57016348]
  [-6.40792615 -6.40794207 -6.41262707 -6.40722306]
  [-6.19563068 -6.19150736 -6.19055137 -6.19008469]
  [-5.94812063 -5.9433751  -5.94397375 -5.94298732]
  [-5.66411456 -5.65894466 -5.65843276 -5.66155133]
  [-5.35078124 -5.34704226 -5.34938056 -5.35235109]
  [-4.99885228 -4.99516469 -4.99846572 -5.00306954]
  [-4.600921   -4.59841477 -4.59874203 -4.5972695 ]
  [-4.152033   -4.15241057 -4.154525   -4.15715972]
  [-3.66646721 -3.66500498 -3.66553048 -3.66373331]
  [-3.15142676 -3.14236339 -3.14232914 -3.14794916]
  [-2.62328346 -2.62218288 -2.6148902  -2.62865687]]

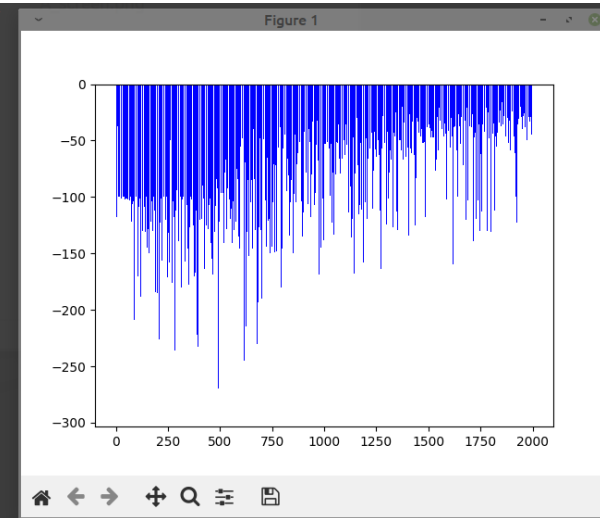
[[ -6.7159586  -6.71529783 -6.71919877 -6.71900845]
 [ -6.47374734 -6.47359043 -6.84136547 -6.47542736]
 [ -6.21253041 -6.21002613 -6.21058457 -6.21021332]
 [ -5.9561723  -5.95173097 -6.02177763 -5.95312993]]
```



Q-Learning

```
laphisboy@laphisboy-900X3K:~/RL_fall/fall_week_1$ python3 cliff_q.py
Success rate: -90.979
Final Q-Table Values
[[[-5.65282542 -5.65361114 -5.65574407 -5.65471412]
  [-5.53672858 -5.5375399  -5.54061073 -5.53960431]
  [-5.35812053 -5.35881434 -5.3613277  -5.36680103]
  [-5.14629029 -5.14456665 -5.14268356 -5.14482495]
  [-4.89721249 -4.89824816 -4.8961145  -4.89803612]
  [-4.62018105 -4.61822491 -4.61625788 -4.62235642]
  [-4.30511984 -4.30426263 -4.30310824 -4.30619001]
  [-3.95317341 -3.95352794 -3.95746675 -3.95710484]
  [-3.56627702 -3.56746228 -3.56728727 -3.57233242]
  [-3.15481121 -3.15248086 -3.15664835 -3.15322673]
  [-2.73066757 -2.72614154 -2.72950243 -2.7296559 ]
  [-2.35850521 -2.35794923 -2.35953948 -2.3586992 ]]]

[[ -5.76232745  -5.76456868  -5.76388702  -5.76406465]
 [ -5.60076494  -5.60036745  -5.60415485  -5.60290239]
 [ -5.39606477  -5.39335139  -5.39800854  -5.39444687]
 [ -5.16188336  -5.15905796  -5.16290388  -5.16118663]]
```



Why SARSA rather than Q-learning?


```
21 # action = [up, right, down, left]
```

SARSA

```
55 Q Values near starting position
56 START
57 [[ -5.65773872 -73.19532831 -6.58676967 -6.09652006]]
58 Above START
59 [[ -5.4619545 -5.60901178 -5.85142828 -5.46444349]
60  [ -4.68087528 -4.90508453 -44.73165228 -4.67900451]]
61 Above above START
62 [[ -5.18131607 -5.18191879 -5.17902335 -5.18519691]
63  [ -4.94702408 -4.94652235 -5.08051739 -4.95036558]]
```

Q-Learning

```
55 Q Values near starting position
56 START
57 [[ -4.47977623 -87.76072562 -4.48266823 -4.47962649]]
58 Above START
59 [[ -4.12097592 -4.12070405 -4.11925462 -4.11693268]
60  [ -3.8334134 -3.83093316 -57.86657778 -3.83874233]]
61 Above above START
62 [[ -3.90748219 3.90242961 -3.90697297 -3.90700077]
63  [ -3.75335178 -3.75719326 -3.75394856 -3.75815594]]
```

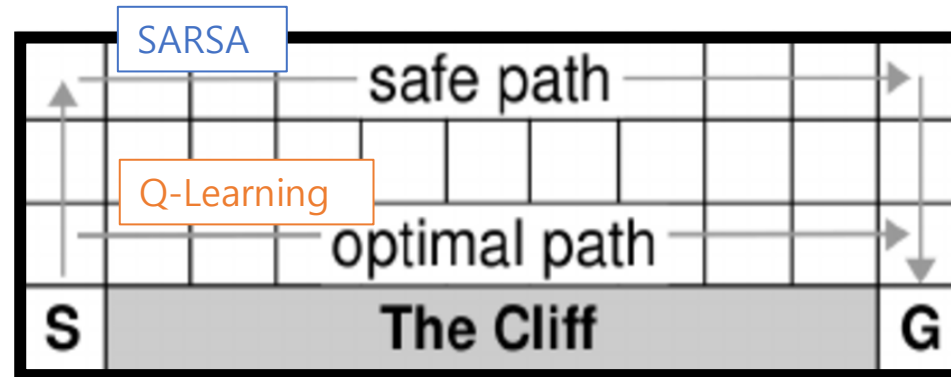


https://github.com/laphisboy/RL_fall/blob/master/fall_week_1/cliff_SARSA_text.txt

https://github.com/laphisboy/RL_fall/blob/master/fall_week_1/cliff_q_text.txt

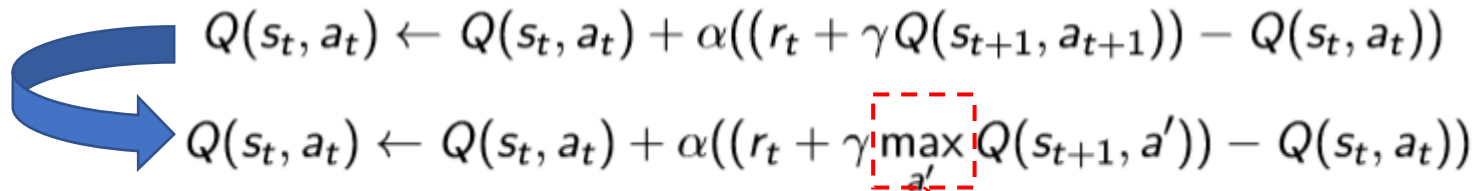
Why SARSA rather than Q-learning?

```
21 # action = [up, right, down, left]
```




R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1
R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1
R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1	R = -1
S	The Cliff (R = -100)								G

SARSA \rightarrow Q-Learning


$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_t + \gamma Q(s_{t+1}, a_{t+1})) - Q(s_t, a_t))$$
$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_t + \gamma \max_{a'} Q(s_{t+1}, a')) - Q(s_t, a_t))$$

Leads to positive bias
= Maximization Bias

Maximization Bias




$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_t + \gamma Q(s_{t+1}, a_{t+1})) - Q(s_t, a_t))$$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_t + \gamma \max_{a'} Q(s_{t+1}, a')) - Q(s_t, a_t))$$

Q. Both have maximization bias right?

Maximization Bias



$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_t + \gamma Q(s_{t+1}, a_{t+1})) - Q(s_t, a_t))$$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_t + \gamma \max_{a'} Q(s_{t+1}, a')) - Q(s_t, a_t))$$

Q. Both have maximization bias right?


A. Yes because policy is updated with action with maximum Q-value = optimism in e-greedy policy

B. But SARSA is less optimistic while evaluation of given action

Q-Learning → Double Q-Learning

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_t + \gamma \max_{a'} Q(s_{t+1}, a')) - Q(s_t, a_t))$$

$$Q_1(S_t, A_t) \leftarrow Q_1(S_t, A_t) + \alpha \left[R_{t+1} + \gamma Q_2(S_{t+1}, \arg \max_a Q_1(S_{t+1}, a)) - Q_1(S_t, A_t) \right]$$



Q7. What is double Q-learning? And how does it overcome maximization bias?


Double Q learning?

```
1: Initialize  $Q_1(s, a)$  and  $Q_2(s, a), \forall s \in S, a \in A$   $t = 0$ , initial state  $s_t = s_0$ 
2: loop
3:   Select  $a_t$  using  $\epsilon$ -greedy  $\pi(s) = \arg \max_a Q_1(s_t, a) + Q_2(s_t, a)$ 
4:   Observe  $(r_t, s_{t+1})$ 
5:   if (with 0.5 probability) then
6:      $Q_1(s_t, a_t) \leftarrow Q_1(s_t, a_t) + \alpha(r_t + \gamma \max_a Q_2(s_{t+1}, a))$ 
7:   else
8:      $Q_2(s_t, a_t) \leftarrow Q_2(s_t, a_t) + \alpha(r_t + \gamma \max_a Q_1(s_{t+1}, a))$ 
9:   end if
10:   $t = t + 1$ 
11: end loop
```

Q-Learning → Double Q-Learning

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_t + \gamma \max_{a'} Q(s_{t+1}, a')) - Q(s_t, a_t))$$

$$Q_1(S_t, A_t) \leftarrow Q_1(S_t, A_t) + \alpha \left[R_{t+1} + \gamma Q_2(S_{t+1}, \arg \max_a Q_1(S_{t+1}, a)) - Q_1(S_t, A_t) \right]$$



Q7. What is double Q-learning? And how does it overcome maximization bias?

한 번 왼쪽 action 에서 0보다 큰 값이 나오면 지속적으로 argmax가 left action이 되는 문제가 생김 (true action(?) = right)

Double Q-learning 은 해결이 아닌 완화 method으로 이해...

참고

Base code <https://github.com/hunkim/ReinforcementZeroToAll>

See how SARSA is implemented <https://github.com/rlcode/reinforcement-learning/tree/master/1-grid-world>

Cliff environment <https://github.com/podondra/gym-gridworlds>

Difference
between those
two https://tcnguyen.github.io/reinforcement_learning/sarsa_vs_q_learning.html

코드

https://github.com/laphisboy/RL_fall/tree/master/fall_week_1